

Dynamic Particle Simulations using GPUs



3 Simulations.

What is a scheduler?

- Launching kernels
- Work distribution
- Parameters
- Coordination

Tests and analysis of schedulers.

Architecture

Specifications of GPU:

- No. Streaming Multiprocessors (SMs): 80
- No. Processing Blocks pr. SM: 4
- Max Warps pr. Processing Blocks: 16
- No. Threads pr. Warp: 32

Maximum No. Threads pr. SM: $4 \times 16 \times 32 = 2048$

Maximum No. Threads: $2048 \times 80 = 163.840$




Scheduler Tests

Simple simulation to analyse the performance of various scheduling techniques

What does the simulation do?

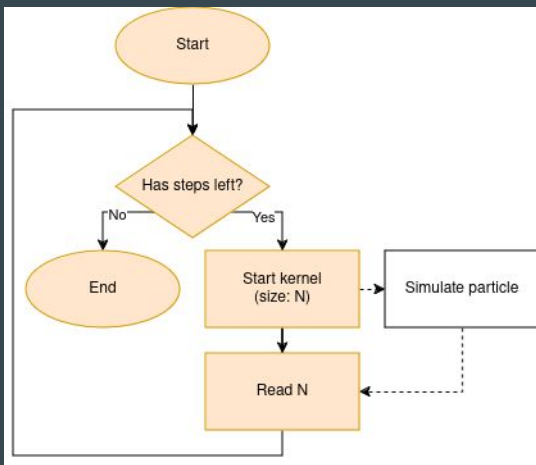
- Moves in a 2D space
- Particles are affected by gravity
- Split when they hit the bottom



Scheduler Tests

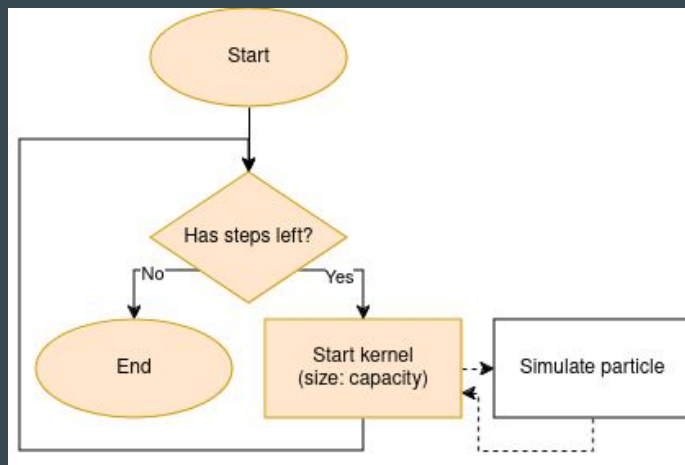
- For all schedulers: Each thread handles one particle

CPU Sync Iterate



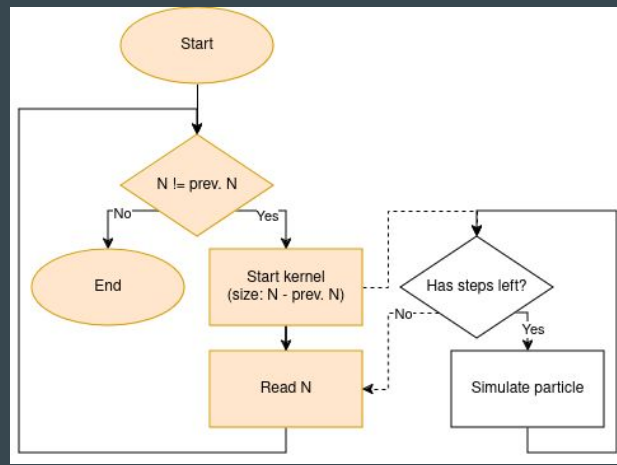
1 thread pr. particle
currently being simulated

Huge Iterate



1 thread pr. possible particle

CPU Sync Full

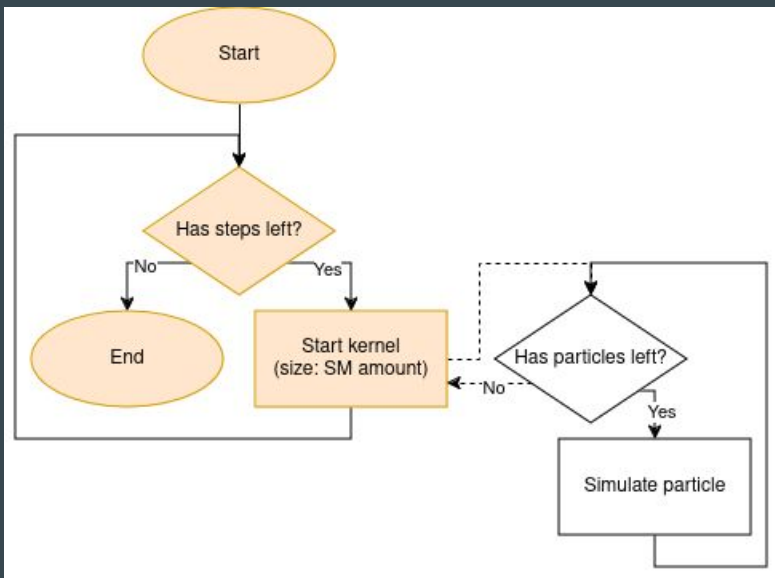


Each thread simulates
particle fully

Scheduler Tests

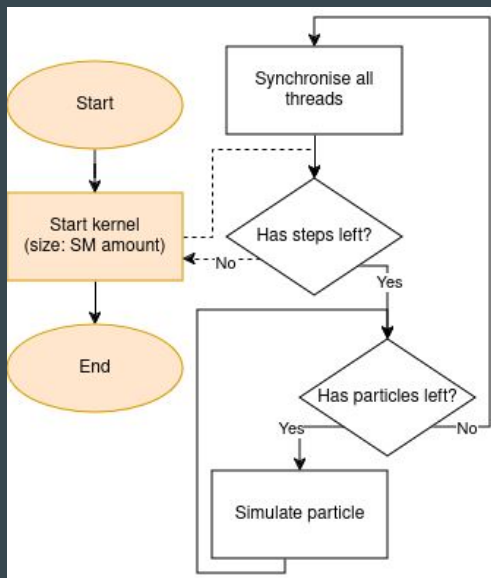
- For all schedulers: Each thread handles multiple particles

Static Simple Iterate



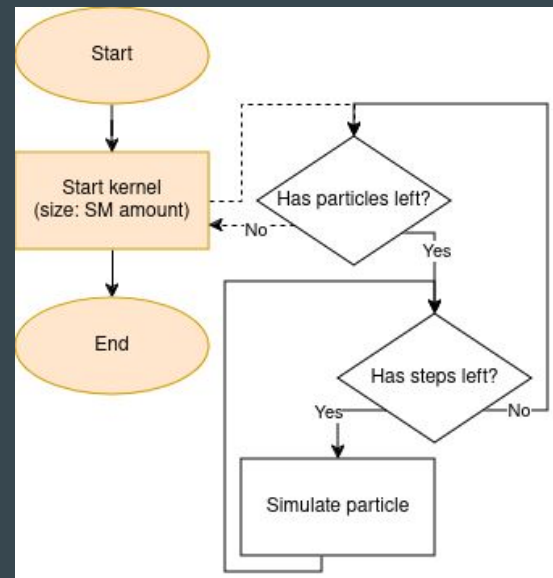
Extension of CPU Sync Iterate

Static GPU Iterate X




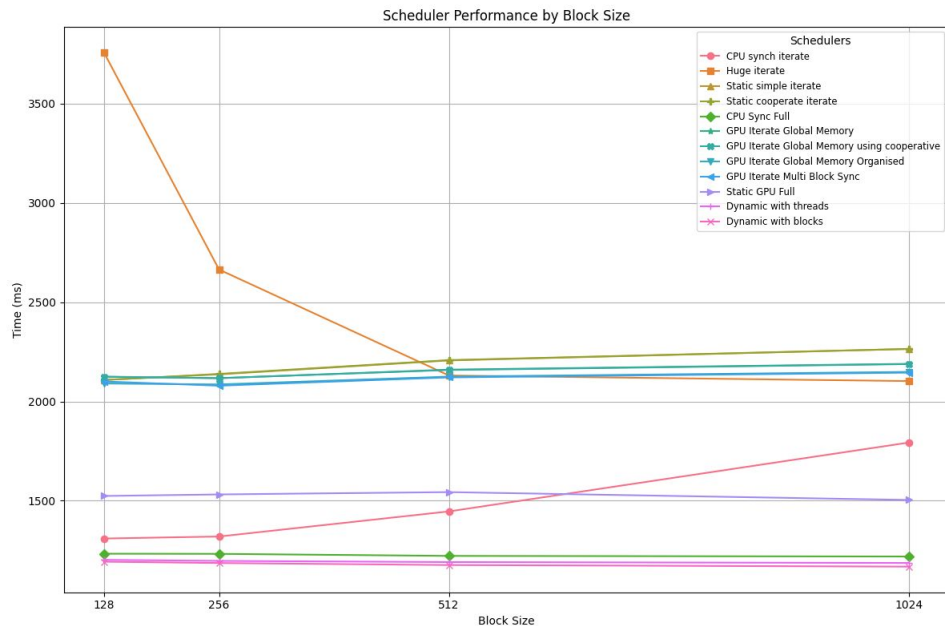
Syncs on GPU

Static GPU Full + Dynamic



Each thread simulates particle fully

Scheduler Tests: Block Size



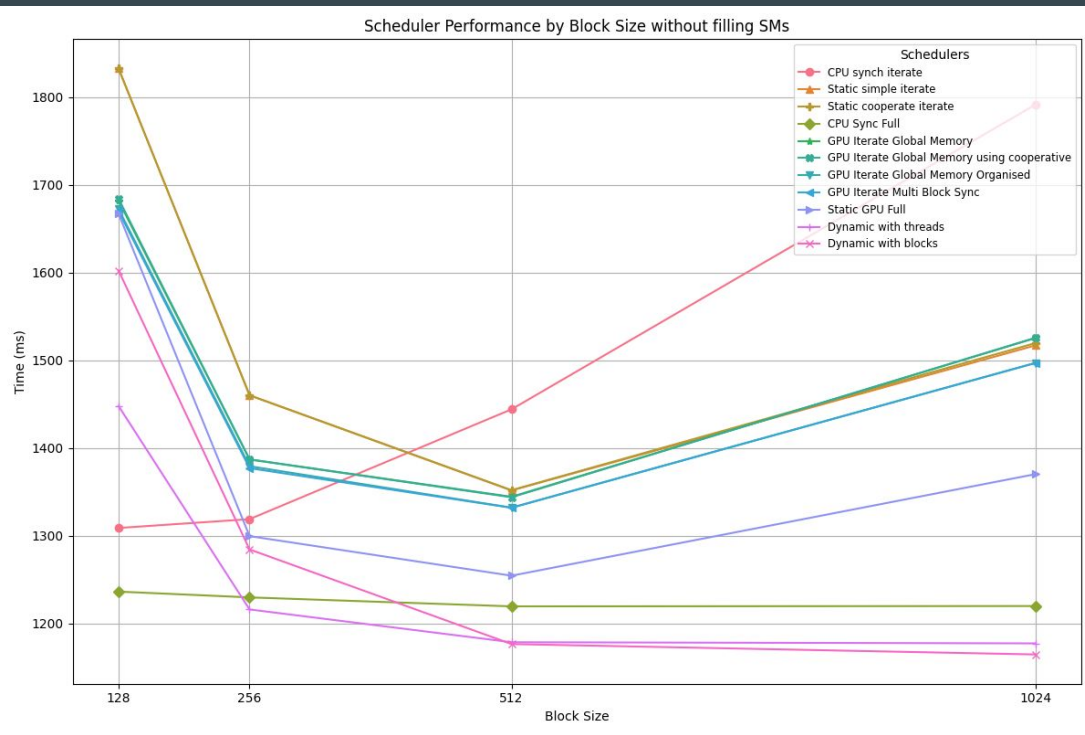
Huge Iterate:

- As many threads as capacity
- More blocks result in more scheduling work

CPU Sync Iterate:

- As many threads as current particles
- Slow threads cause a more severe bottleneck

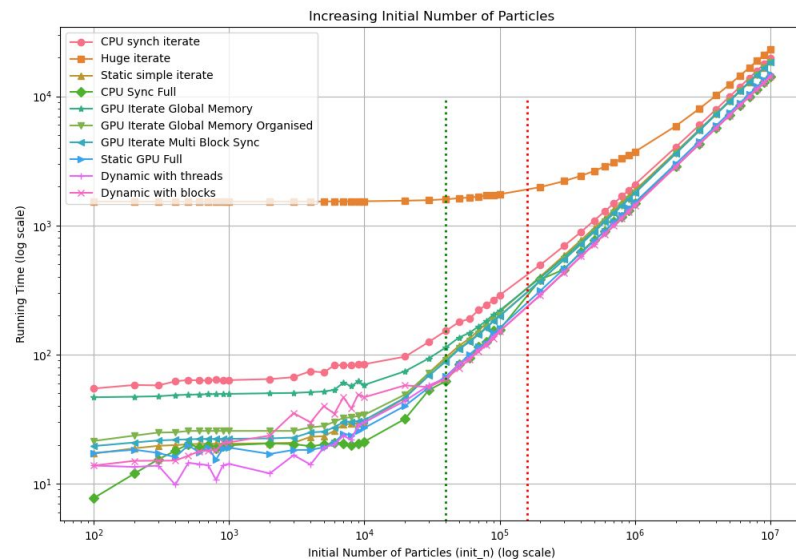
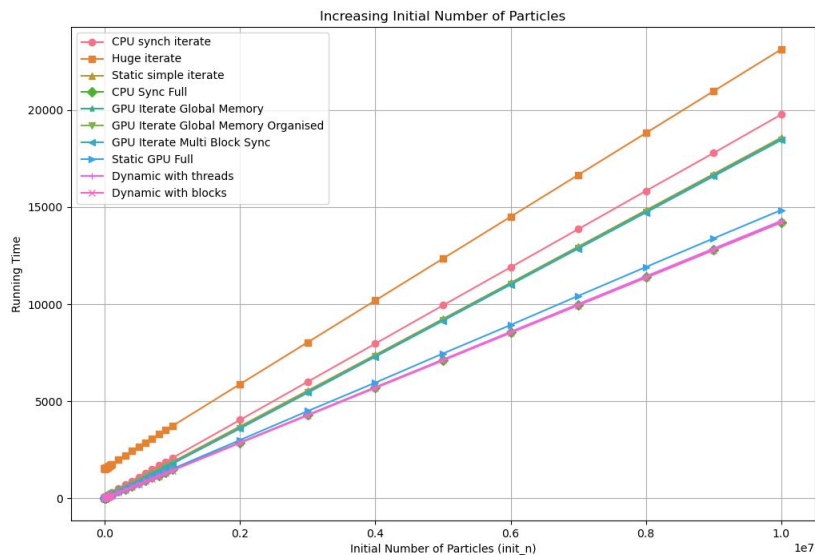
Scheduler Tests: Block Size without filling SMs



Improvement of the Statically Implemented Iterate Schedulers:

- Fewer threads with more work > more threads with less work
- All threads synchronise after each time step
- Extra scheduling work

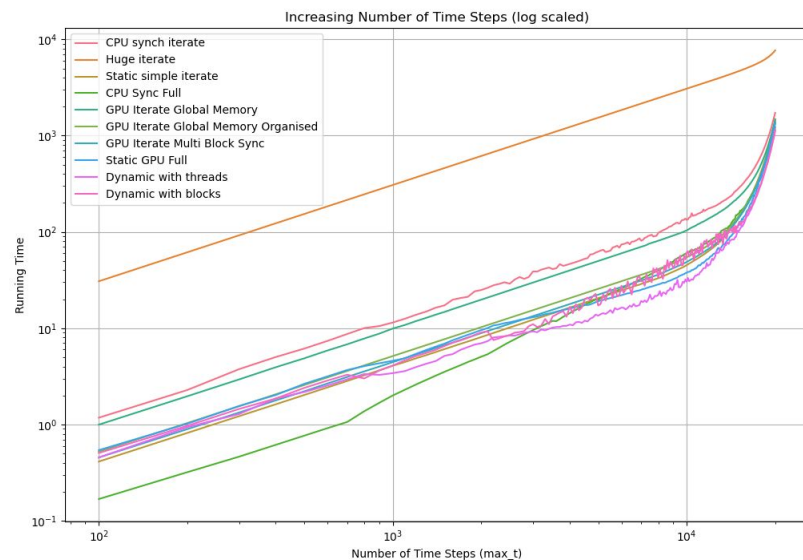
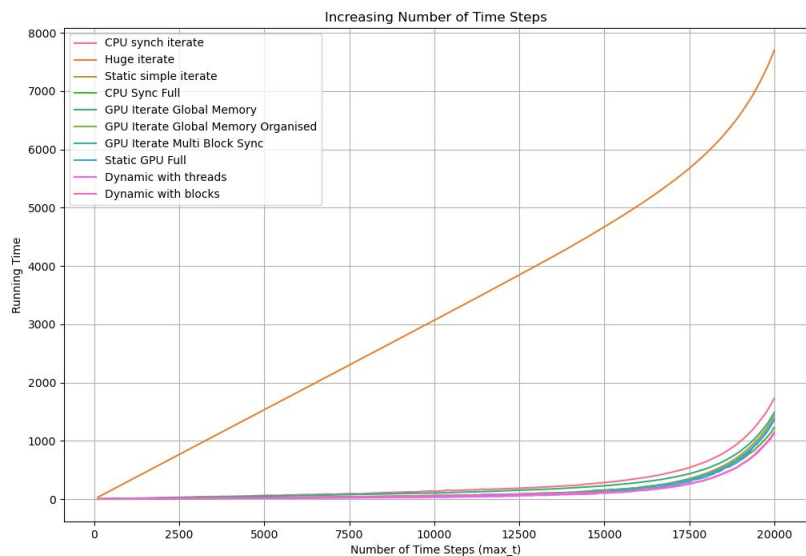
Scheduler Tests: Initial Number of Particles



Linear Growth

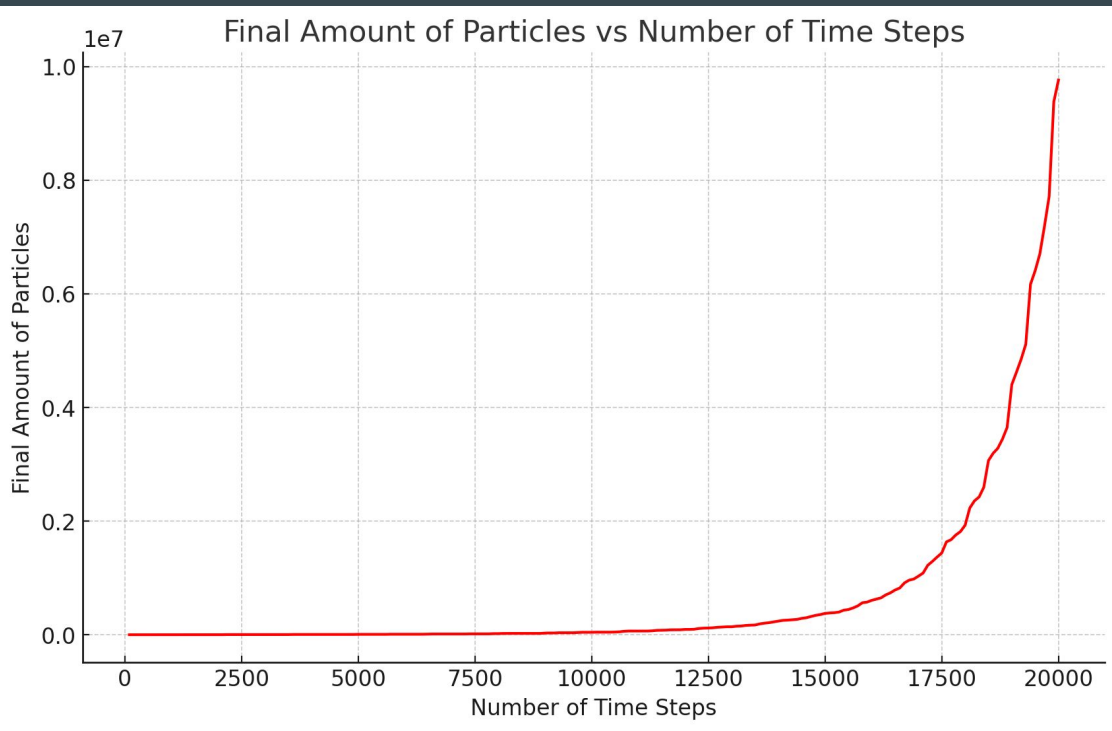
Constant until SMs are filled

Scheduler Tests: Number of Time Steps



Linear until exponential
Exponential when SMs are filled

Scheduler Tests: Number of Time Steps



Exponential growth in number of particles

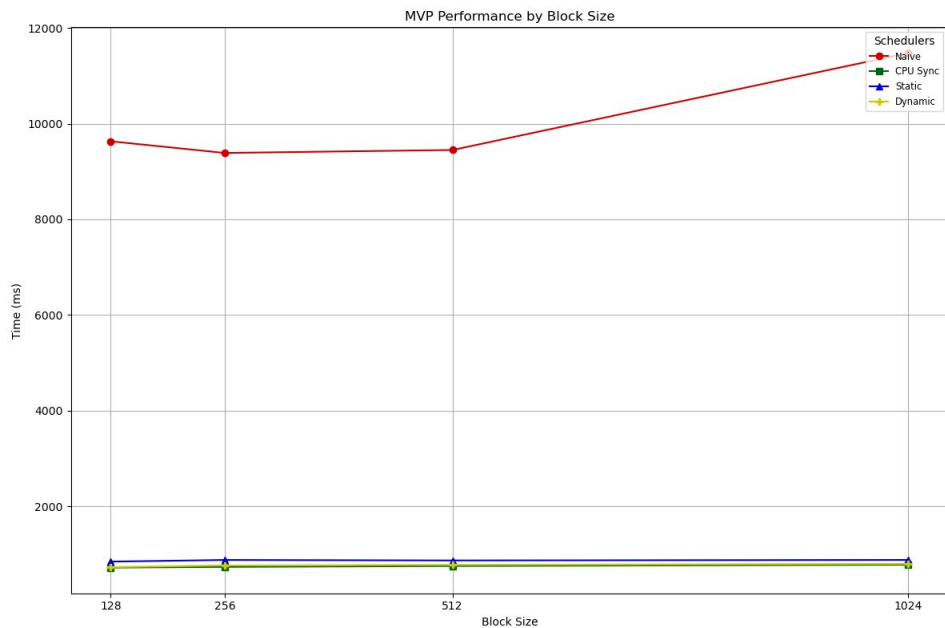
Minimal Viable Product (MVP)

Naive	Based on CPU Sync Iterate. Optimisations: Shared memory and memory organising.
CPU Sync	Same as CPU Sync Full.
Static	Based on Static GPU Full. Optimisations: fewer atomicAdd-operations.
Dynamic	Based on Dynamic With Blocks. Optimisations: Shared memory


Simulation changes:

- Chance to split every step instead of on bounce.
- Changes to particle movement.

MVP: Block Size



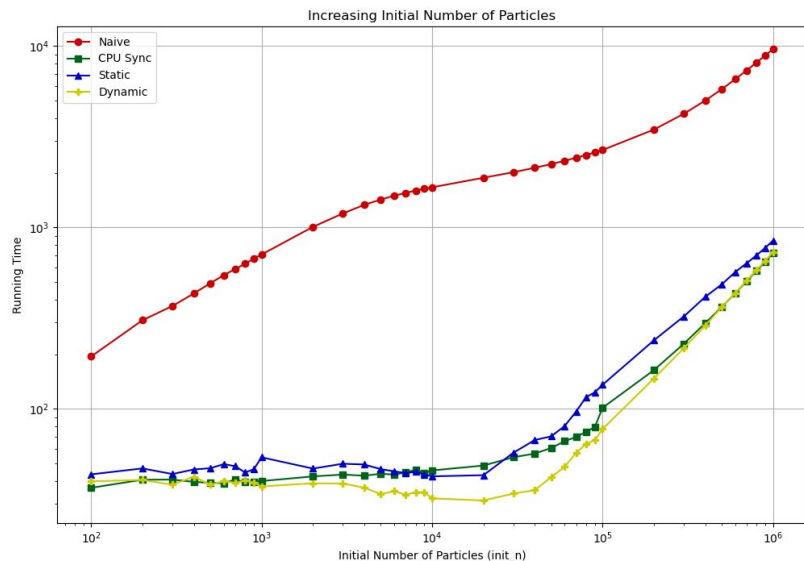
Same as before, but memory optimisation skewed the optimal block size



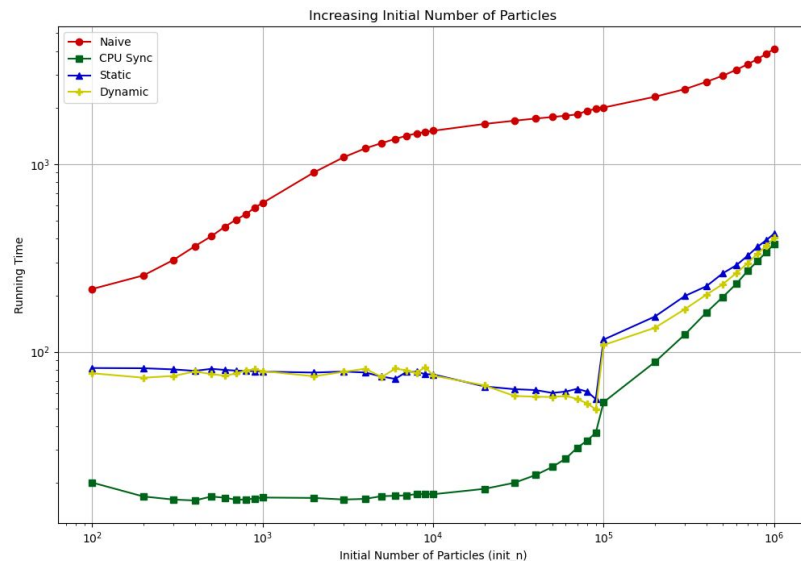
Small block size results in less bottleneck from slow threads

MVP: Initial Number of Particles

Simulating newly spawned



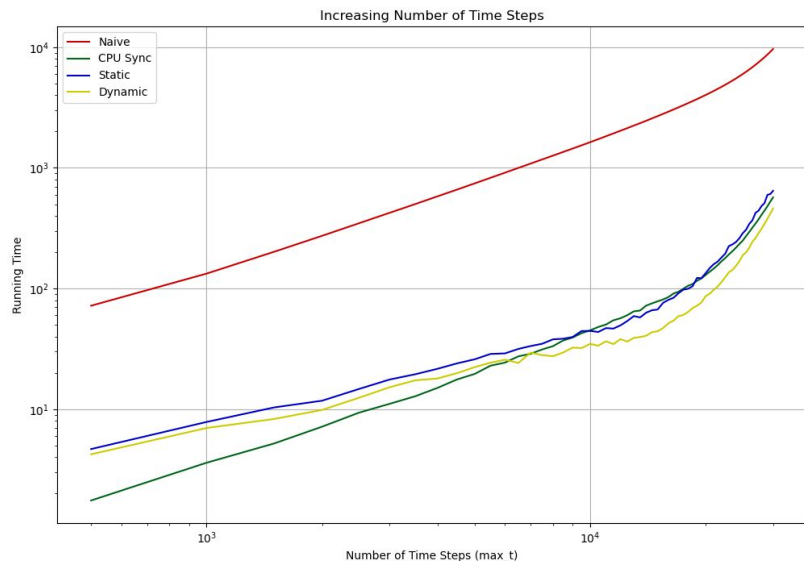
Without simulating newly spawned




for $\text{init_n} > 10.000$
 $\text{init_n} \cdot 0,05 \% > 0,5$

MVP: Number of Time Steps

Simulating newly spawned



Without simulating newly spawned



Exponential growth is caused by increased no. particles

Scheduler For Further Work

Dynamic:

- For the most part performed the best
 - Close behind when not
- Room for optimisations

Static:

- Has many of the same same quality
- But overall worse

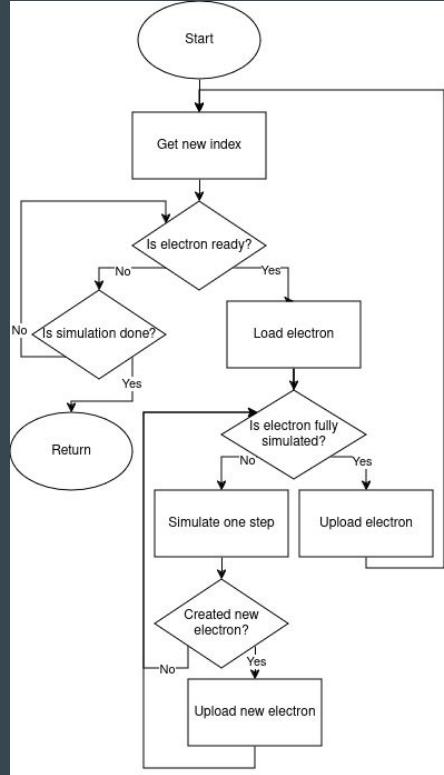
Particle-In-Cell (PIC)

- The final simulation
- Close to physically correct
 - Electrons divided into cells on a grid
 - Poisson steps, each containing many mobility steps
 - Chance for removal of particles
 - Cross sections for collision chances depending on velocity
- Optimises Dynamic

Dynamic Optimisation: Implementation

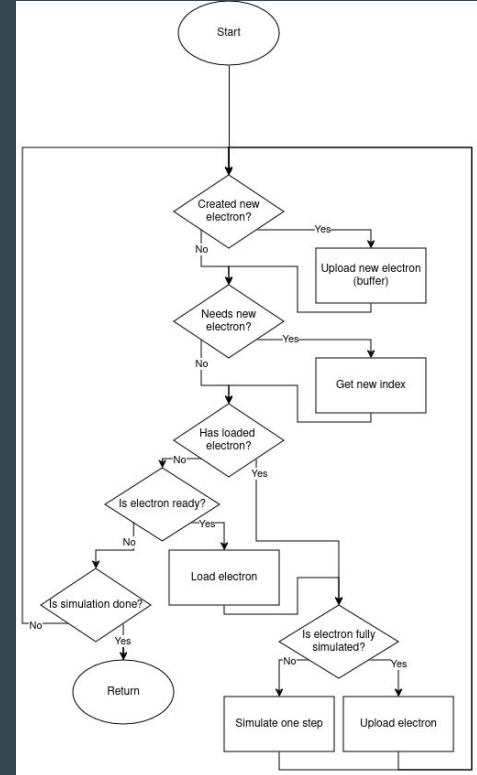
Old:

- Nested loops
- Simple



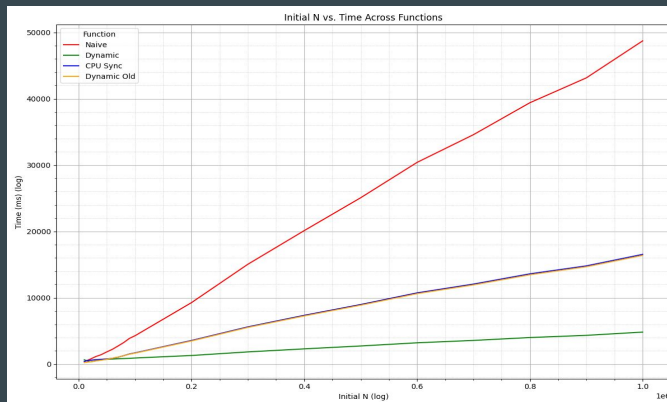
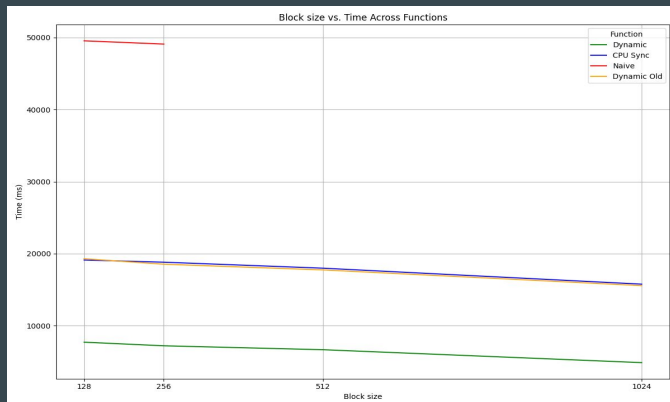
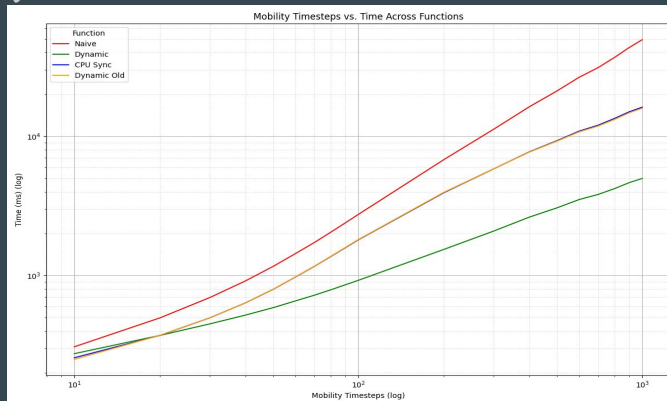
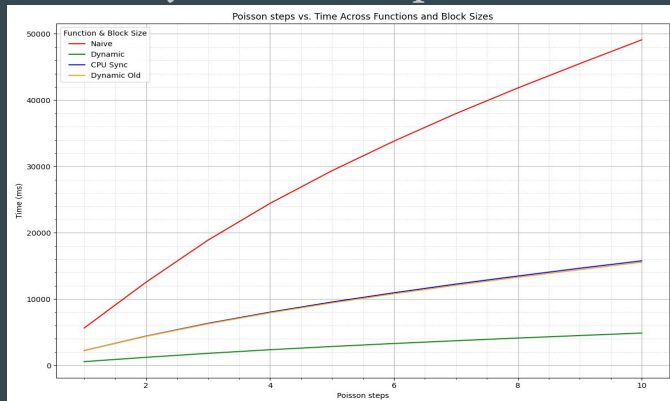
New:

- Only one loop
- More steps
- Allows electron buffer upload
- Warp cooperation



Dynamic Optimisation: Performance

New Dynamic is superior in nearly all cases.



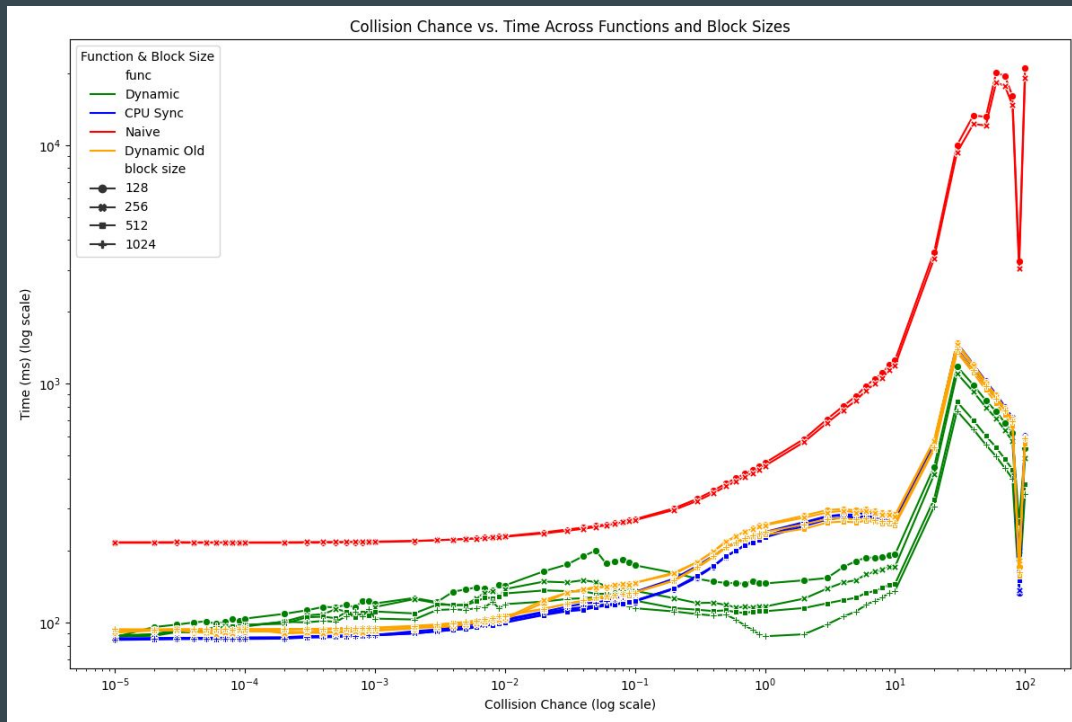
Except...

Dynamic Optimised: Performance for Constant CC

When collision chance is constant:


- If collisions are rare, the additions are not worth it.
- If collisions are very common, the additions make little difference.

But something here is weird.




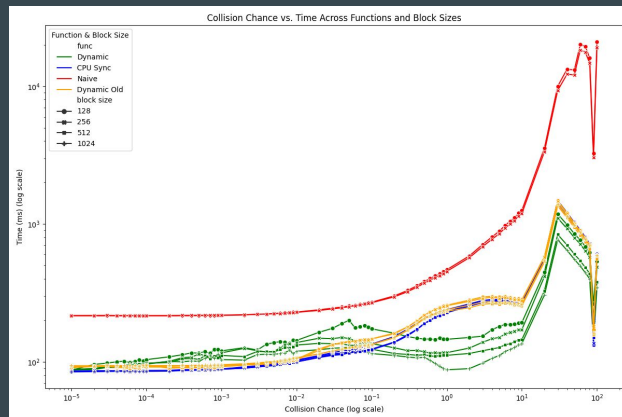
Random Numbers: Weird Behaviour

- Decline after ~30%
- Dip at 90%



Random Numbers: Non-Linear Collision Counts

- The same tendency can be seen in collision counts.
- Caused by lack of uniformity in random numbers.
- Too fast increase after 10% (hits capacity).
 - This explains steep rise in processing time.
- Dip at 90%.



Random Numbers: Use Sequence Parameter

Submitted code:

- Varying seed
- Constant sequence

```
_device_ void newRandState(curandState* rand_state, int seed){  
    curand_init(39587 + seed, 0, 0, rand_state);  
}
```

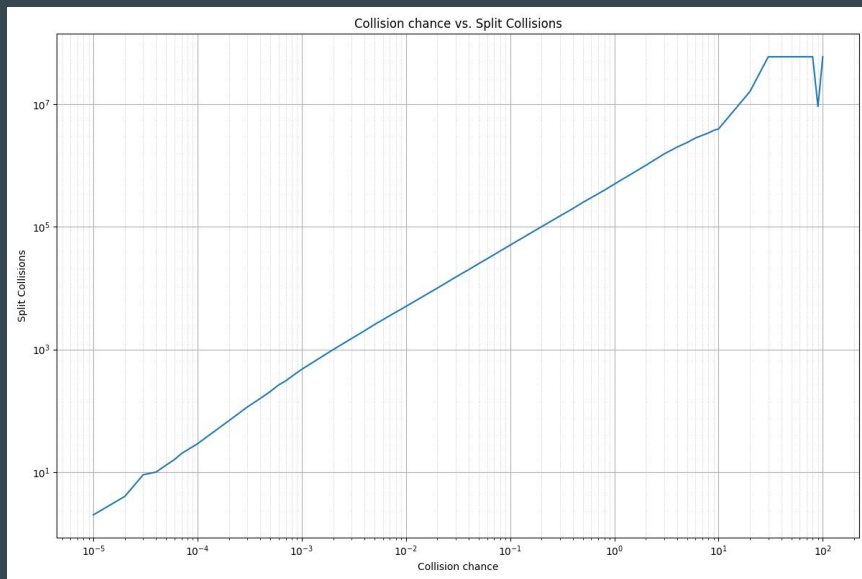
Fixed code:

- Constant seed
- Varying sequence


```
_device_ void newRandState(curandState* rand_state, int sequence){  
    curand_init(39587, sequence, 0, rand_state);  
}
```


Random Numbers: Sequence Fix Results

Before



After



Random Numbers: No Determinism

Submitted code:

- Each particle has a state
- New states generated with new particles

```
d_electrons[new_i] = new_electron;  
d_rand_states[new_i] = new_rand_state;
```

```
newRandState(new_rand_state, randInt(rand_state, 0, 2000000000)+1000000000);
```


Fixed code:

- States initialised at start
- No control of which electron uses which state


```
__global__ void setup_rand(curandState* d_rand_states) {  
    int i = threadIdx.x+blockDim.x*blockIdx.x;  
    newRandState(&d_rand_states[i], i);  
}
```

Random Numbers: No Determinism Results (Collisions)


Before



Varying sequence






No determinism






Random Numbers: Time Results

L:



S:




Before

Varying sequence


No determinism

Random Numbers: Consequences for Mobility Steps


Before



Varying sequence



No determinism



Conclusion

- We have come to understand various scheduling paradigms.
 - How they influence performance.
 - Why they work as they do.
 - How the details of the simulation affects the performance of various schedulers.
 - Persistent dynamic schedulers have a lot of potential and are often superior.
- We have created a near-realistic simulation and optimised a scheduler for it.
- Random number generation matters.
 - Can affect outcome.
 - Can affect benchmarks.
- Beyond schedulers, impact of changes can be hard to predict.
 - Always test different block sizes.
 - Micro-benchmarks are helpful for small changes.